


# Parallelism

## w2


Learning to communicate

# MPI - What we have learned


MPI PointToPoint  
out: MPI\_RANK



MPI Point2Point  
out: MPI\_RANK



# MPI\_Send()



# MPI\_Send()

## MPI\_Send

Performs a blocking send

### Synopsis

```
int MPI_Send(const void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)
```

### Input Parameters

**buf**  
initial address of send buffer (choice)

**count**  
number of elements in send buffer (nonnegative integer)

**datatype**  
datatype of each send buffer element (handle)

**dest**  
rank of destination (integer)


**tag**  
message tag (integer)

**comm**  
communicator (handle)

### Notes

This routine may block until the message is received by the destination process.

# MPI\_Recv()



# MPI\_Recv()

## MPI\_Recv

Blocking receive for a message

### Synopsis

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag,  
             MPI_Comm comm, MPI_Status *status)
```

### Output Parameters

**buf**  
initial address of receive buffer (choice)

**status**  
status object (Status)

### Input Parameters

**count**  
maximum number of elements in receive buffer (integer)

**datatype**  
datatype of each receive buffer element (handle)


**source**  
rank of source (integer)

**tag**  
message tag (integer)

**comm**  
communicator (handle)

# Blocking communication


- This aliasing of send() and recv() is to done to avoid deadlocks.
- Both MPI\_Send() and MPI\_Recv() are blocking. **(They wait for completion)**
- Use “tags” to pair them
- These scopes are parallel threads



```
...  
  
MPI_Status stat;  
  
// Allocate memory for A,B on CPU  
vector<double> A_vect(N), B_vect(N);  
double* A = A_vect.data();  
double* B = B_vect.data();  
  
int tag1 = 10; // can be any unique integer  
int tag2 = 20; // can be any unique integer  
  
for(int i=1; i<=5; i++){  
    1 if(rank == 0){  
        2 MPI_Send(A, N, MPI_DOUBLE, 1, tag1, MPI_COMM_WORLD);  
        MPI_Recv(B, N, MPI_DOUBLE, 1, tag2, MPI_COMM_WORLD, &stat);  
    }  
  
    1 else if(rank == 1){  
        MPI_Recv(B, N, MPI_DOUBLE, 0, tag1, MPI_COMM_WORLD, &stat);  
        2 MPI_Send(A, N, MPI_DOUBLE, 0, tag2, MPI_COMM_WORLD);  
    }  
}  
  
...
```


15	12
1	1

# Application: Sequential Broadcast





# Application: Ring Sequential



Must be aware of deadlocks:

1. send from even to odd
2. send from odd to even

# HyperCube Broadcast

- Hypercube broadcasting is a communication pattern often used in parallel computing, where data needs to be distributed from one process (typically the root) to all other processes in a structured manner.
- In the context of MPI (Message Passing Interface), you can perform hypercube broadcasting using `MPI_Send` and `MPI_Recv` calls.

